

Introduction to OpenMP

Thomas Hauser
Director of Research Computing
University of Colorado Boulder
[USGS Alaska Survey](#)
08.26.2014



Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Research Computing @ CU Boulder

USGS parallel computing workshop

3/14/17

How to Run Applications Faster?

- There are 3 ways to improve performance:
 - Work Harder
 - Work Smarter
 - Get Help
- Computer Analogy
 - Using faster hardware
 - Optimized algorithms and techniques used to solve computational tasks
 - Multiple computers to solve a particular task

Research Computing @ CU Boulder

USGS parallel computing workshop

3/14/17

Parallelism on a Processor

- Scaling
 - Using more cores
 - Task parallelism
- Vectorization
 - Wider vectors (same operation on larger number of data)
 - Data parallelism

Research Computing @ CU Boulder

USGS parallel computing workshop

3/14/17

Parallelism

• Nodes – MPI

- The following picture shows four stacks aligned side-by-side.



• Threads – OpenMP

- The following pictures shows ...




• Instructions – ILP

- The following is three lines of instruction code

```
I1: add R1, R2, R3
I2: sub R4, R1, R5
I3: xor R10, R2, R11
```

• Data – SIMD

- The following picture shows ...



Research Computing @ CU Boulder

USGS parallel computing workshop

3/14/17

OpenMP

- The OpenMP application programming interface (API) supports multi-platform shared-memory parallel programming in
 - C/C++
 - Fortran
- The API defines a portable, scalable model with a simple and flexible interface for developing parallel applications on platforms from the desktop to the supercomputer.

[OpenMP Website](#)

Philosophy

- Goal: Add parallelism to a functioning serial code.
- Requires: Shared memory machine.
- How: Add *compiler directives* to parallelize parts of code.
- Pro: Often very easy to add to existing code.
- Con: Large shared memory machines are expensive.

Incremental Parallelization

- Sequential program a special case of a shared-memory parallel program
- Parallel shared-memory programs may only have a single parallel loop
- Incremental parallelization: process of converting a sequential program to a parallel program a little bit at a time

Resources

- *Using OpenMP* the book
 - [The Mit Press Website](#)
- API Quick Reference
 - [OpenMP C & C++ PDF Document](#)
 - [OpenMP Fortran PDF Document](#)
- OpenMP v4 full API
 - [OpenMP 4.0.0 PDF Document](#)
- OpenMP Examples
 - [OpenMP 4.0.2 PDF Document](#)


What's OpenMP Good For?

- C/Fortran + OpenMP sufficient to program shared memory computers
- C/Fortran + MPI + OpenMP a good way to program multicomputers built out of multiprocessors
 - Cluster nodes
 - Intel Phi accelerators

Fork/Join Parallelism

- Program starts as a single thread of execution, initial thread.
- A team of threads is forked at the beginning of a parallel region.
- At the end of a parallel region the threads join (either die or are suspended).

• The following diagram shows the flow between the initial thread through the fork, to the team of threads to join and become the initial thread.




Memory Model

- Contents of memory segments:
 - static variables
 - variables on the run-time stack
 - functions on the run-time stack
 - dynamically allocated data on the heap
- Below is a snippet of code designed to ...

```
#include <stdlib.h> static
const int n = 4; int
main(int argc, char **argv)
{
    int i = 0;
    int *c = NULL;
    c = malloc(n * sizeof(int));
    for (i = 0; i < n; ++i) {
        c[i] = i + 1;
    }
    free (c);
    return(0);
}
```

The following diagram demonstrates Program Memory



Compiler Directive

- Compiler directive in Fortran
- A way for the programmer to communicate with the compiler
- Compiler free to ignore directives
- Syntax:


```
!$OMP <rest of directive>
!$OMP& continuation line
!$ INTEGER :: i !something only in the
parallel version of the program
```

Pragma in C or C++

- Syntax:

```
#pragma omp <rest of directive>
```

Compiler Directives

This table provides the following information: ...

| Parallel Control | Data | Work Sharing | Synchronization | Scheduling |
|---|--|---|--|---|
| Controls the flow of parallel regions parallel | Specifies variables scope shared private | Distribution of work between threads for do | Coordination of threads critical atomic barrier | Loop iteration distribution schedule |
| Vectorization | | Accelerators | | |
| Loop vectorization simd | | Offload to Co-processors GPUs target | | |

Using the GCC compilers with OpenMP

- `gfortran -fopenmp -g basic.f90`
- `gcc -fopenmp -g basic.c`
- `setenv OMP_NUM_THREADS 12` if `tcsh/csh`
- `export OMP_NUM_THREADS=12` if `bash/sh`
- `./a.out`

Using the Intel compilers with OpenMP

- `ifort -fopenmp -g basic.f90`
- `icc -fopenmp -g basic.c`
- `setenv OMP_NUM_THREADS 12` if `tcsh/csh`
- `export OMP_NUM_THREADS=12` if `bash/sh`
- `./a.out`

Parallel

Compiler must be able to verify the run-time system will have information it needs to schedule loop iterations

Fortran

```
!$OMP parallel do
do i=1,10
    a(i) = b(i) + c(i)
end do
!$OMP end parallel do
```

C

```
#pragma omp for
for (i=0;i<10;i++){
    a[i] = b[i] + c[i];
}
```

OpenMP – Workshare Directives

```
!$OMP PARALLEL [clause]
Tid = omp_get_thread_num()
!$OMP DO
    Do l = 1, nmax
    End do

!$OMP END DO
!$OMP END PARALLEL [nowait]
```

parallel directive

- The `parallel` directive precedes a block of code that should be executed by *all* of the threads
- Note: execution is replicated among all threads

do directive

- The `parallel` directive instructs every thread to execute all of the code inside the block
- If we encounter a `do` loop that we want to divide among threads, we use the `do` directive

```
!$OMP do
```

Example use of do Directive

The code below is an example of do Directive. The lines beginning with !\$OMP are presented in red text.

```

!$omp parallel private(i,j)
Do i = 1, m
  low = a(i)
  high = b(i)
  if (low > high) then
    write(*,*) "Exiting", i
    exit
  endif
End do
!$omp do
do j = low, high
  c(j) = (c(j) - a(i))/b(i)
End do
!$omp end do
!$omp end parallel

```

Parallel Regions

- We tell OpenMP compiler to parallelize code.
- Mark parallel blocks.
- The compiler will spawn threads and split the work up.
- We can tell the compiler the number of threads too.

- Parallel code for C:

```

#pragma omp
parallel
{
    ...
}

```

- Parallel code for Fortran:

```

!$OMP parallel
...
!$OMP end parallel

```

OpenMP Also Provides Library Calls.

C function prototypes are in `omp.h`.

Fortran module interface is in `omp_lib`.

For compatibility, you should `#ifdef` guard these calls.

Remember to use the pre-processor for Fortran too (. F90).

- Code for C:

```
#ifdef _OpenMP
#include <omp.h>
#endif
```

- Code for Fortran:

```
#ifdef _OpenMP
  use omp_lib
#endif
```

Library Functions

- OpenMP has functions to find out our thread number and the total number of threads.
 - `omp_get_thread_num()`
 - `omp_get_num_threads()`
- We can set the number of threads:
 - Function `omp_set_thread_num()`.
 - Specifying `num_threads` clause after the `parallel` directive.
 - With the environment variable **OMP_NUM_THREADS**.

nthreads/nthreads_c.c

```

#include <stdlib.h>
#include <stdio.h>
#include <omp.h>

int
main(int argc, char **argv)
{
    #pragma omp parallel
    {
        printf("Hello world! From thread %d\n", omp_get_thread_num());
    } /* End omp parallel */

    return(EXIT_SUCCESS);
}

```

There is also a Fortran version nthreads/nthreads_f.f90.

Compiling the program

```

GCC gcc -fopenmp -o thread_num_c thread_num_c.c
Intel icc -openmp -o thread_num_c thread_num_c.c
IBM xlc -qsmp=omp -o thread_num_c thread_num_c.c

```

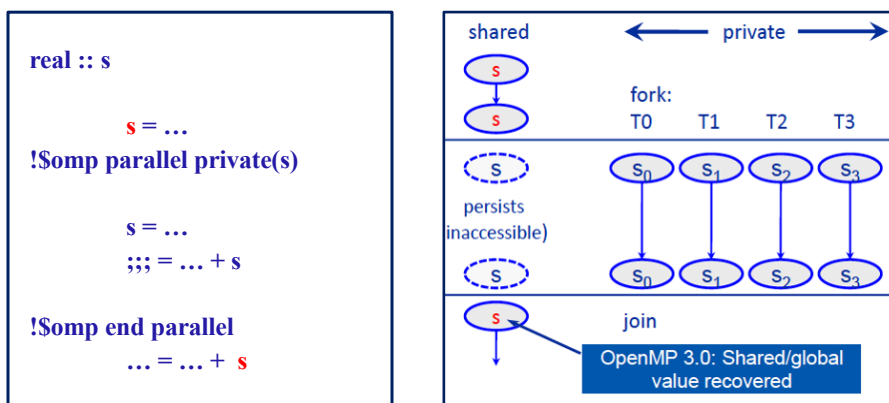
- Execute the program, specifying different numbers of threads.
 1. ./thread_num_c
 2. env OMP_NUM_THREADS=1 ./thread_num_c
 3. env OMP_NUM_THREADS=24 ./thread_num_c
- What is the output?
 - Threads printed out their identification number.
 - Random order of numbers. Threads execute independently and in general order will be random.

Data Scoping – shared vs private

- Default: All data in a parallel region is shared
- This includes global data (global/static variables, C++ class variables)
- Exceptions:
 - Local data within enclosed function calls are private (Note: Inlining must be treated correctly by compiler!) unless declared static)
 - Loop variables of parallel (“sliced”) loops are private (cf. workshare constructs)
- Due to stack size limits it may be necessary to make large arrays static
 - This presupposes it is safe to do so!
 - If not: make data dynamically allocated
 - As of OpenMP 3.0: OMP_STACKSIZE may be set at run time (increase thread-specific stack size):
\$ setenv OMP_STACKSIZE 100M

Private Variables – Masking

The following diagrams...



Declaring Private Variables

The following code is ...

```
do j = 1,nj
  do i = 1, ni
    a(i,j) = min(a(i,j), a(i,j)+tmp)
  end do
end do
```

- Either loop could be executed in parallel
- We prefer to make outer loop parallel, to reduce number of forks/joins
- We then must give each thread its own private copy of variable *i*

private Clause

- Clause: an optional, additional component to a pragma
- Private clause: directs compiler to make one or more variables private

```
private ( <variable list> )
```

Example Use of private Clause

```
!$OMP parallel do private(i)
DO j = 1, nj
  DO i = 1, ni
    a(i,j) = MIN(a(i,j),a(i,k)+tmp)
  END DO
END DO
```

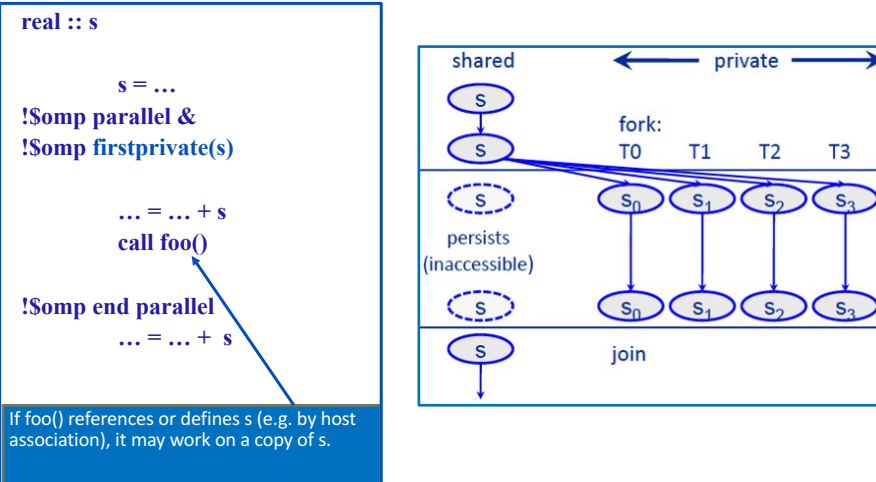
```
#pragma omp for private(i)
for (j=0;j<nj;j++)
  for (i=0; i<ni; i++) {
    a[j,i] = a[j,i]+tmp
  }
```

firstprivate Clause

- Used to create private variables having initial values identical to the variable controlled by the master thread as the loop is entered
- Variables are initialized once per thread, not once per loop iteration
- If a thread modifies a variable's value in an iteration, subsequent iterations will get the modified value

Scoping – firstprivate

The following diagrams show...



Research Computing @ CU Boulder

USGS parallel computing workshop

3/14/17

Use of firstprivate

```

program wrong
I = 10
!$OMP PARALLEL
PRIVATE (I)
I= I + 1
print *, I
!$OMP END PARALLEL
        
```

```

program correct
I = 10
!$OMP PARALLEL
FIRSTPRIVATE (I)
I= I + 1
print *, I
!$OMP END PARALLEL
        
```

Research Computing @ CU Boulder

USGS parallel computing workshop

3/14/17

lastprivate Clause

- Sequentially last iteration: iteration that occurs last when the loop is executed sequentially
- **lastprivate** clause: used to copy back to the master thread's copy of a variable the private copy of the variable from the thread that executed the sequentially last iteration

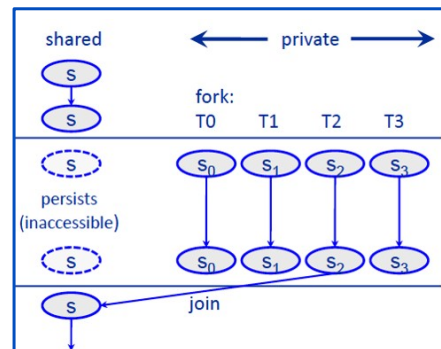
Scoping – lastprivate

The following diagram shows...

```

real :: s

s = ...
!$omp parallel &
!$omp lastprivate(s)
!$omp do
  do i = ...
    s = ...
  end do
!$omp end do
!$omp end parallel
... = ... + s
  
```



Function `omp_get_num_procs`

- Returns number of physical processors available for use by the parallel program
- C: **`int omp_get_num_procs(void);`**
- Fortran:

interface

```
function omp_get_num_procs ()
  use omp_lib_kinds
  integer ( kind=omp_integer_kind ) :: omp_get_num_procs
end function omp_get_num_procs
```

end interface

Function `omp_set_num_threads`

- Uses the parameter value to set the number of threads to be active in parallel sections of code
- May be called at multiple points in a program

```
subroutine omp_set_num_threads (t)
  integer, intent(in) :: t
```

```
void omp_set_num_threads(int num_threads);
```

Variables

- We must tell the compiler how to use variables.
 - A `shared` variable has the same address in memory in every thread.
 - A `private` variable has a different address in memory in every thread.
 - A `first private` is private, however it is pre-initialize.
- Clauses specifies the scope of variables.

Variables in C

```
#pragma omp parallel \  
    default(none) \  
    shared(x) \  
    private(i,j) \
```

Variables in Fortran

```
!$OMP parallel      &  
!$OMP default(none) &  
!$OMP shared(x)    &  
!$OMP private(i,j)  
...  
!$OMP end parallel
```

The default is shared.

- Try and always specify `default(none)`, so as not to confuse a variables behavior. Then explicitly define every variable.
- ► In C, you are able to declare variables within structured blocks to reduce it's scope. This will make it a private variable. For example, `i` is shared while `j` is private.

Default in C

```
int i = 0;  
#pragma omp parallel \  
    default(none) \  
    shared(i)  
{  
  
    int j = 0;  
  
}
```

Questions?

[Online Survey](#)

[Thomas Hauser's Email](#)

License

This work is licensed under the Creative Commons Attribution 4.0 International License. To view a copy of this license, visit [The Creative Commons Website](#)

When attributing this work, please use the following text: "OpenMP", Research Computing, University of Colorado Boulder, 2016. Available under a Creative Commons Attribution 4.0 International License.

The following graphic is the logo for Creative Commons.

